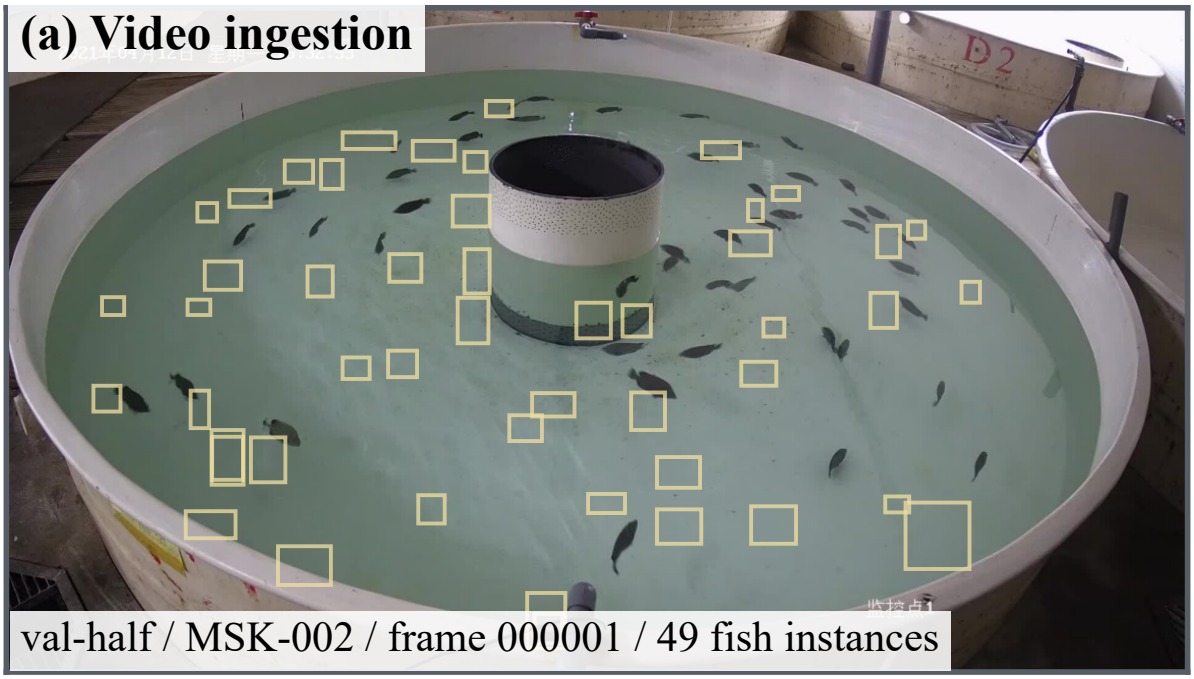

(b) Closed-loop reviewer map

C1

Audit and reproducibility

split=val_half, max_frames=1000
seeds=[0, 1, 2]
run_config.json + dual manifests + release bundle

↓

C2

Stress and stratified diagnostics

global metrics + bucket metrics + stress groups
occlusion / density / turning / low-confidence
localize risk before deployment selection

↓

C3

Window-level drift monitoring

compute D_in(t) and D_out(t) over sliding windows
selected rule: K=1, q=99
false-alert=2.1% / missed-drift=27.1%

↓

OUT

Deployment outputs

HOTA / IDF1 / IDSW / CountMAE
FPS / peak memory / normalized CPU
Base vs +gating policy recommendation